

Mike Leske
R00183658

COMP9058 - Metaheuristic Optimization
Assignment 2 – Local Search

Table of Contents

<i>Preface</i>	3
<i>1 GWSAT</i>	4
1.1 Evaluation of “wp”	5
1.2 Evaluation of “restart” and “max_iterations”	9
1.3 Evaluating instances uf20-01 and uf20-02	11
1.4 Conclusion	11
1.5 Debug GWSAT	12
<i>2 WalkSAT/SKC</i>	13
2.1 Evaluation of “p”	14
2.2 Evaluation of “restart” and “max_iterations”	17
2.3 Evaluation of “tl”	19
2.4 Evaluating instances uf20-01 and uf20-02	19
2.5 Conclusion	22
2.6 Debug WalkSAT	23
<i>3 Run-Time Distribution Evaluation</i>	25
<i>Appendix A</i>	29
<i>References</i>	31

Preface

All experiments documented as part of this report were executed on the same computer:

- Intel(R) Core(TM) i7-9750H CPU @ 2.60GHz
- 16 GB 2400 MHz DDR4

For each experiment the following statistics will be listed:

- Executions Overview: Total & Mean execution runtime in μsec , success rate, total iterations across all executions
- Restarts: Statistics on search restarts
- Total Iterations: Statistics on total search iterations
- Iterations to solve: Statistics for the search runs that solved the instance

Note: Statistics shown exclude unsuccessful executions according to lecture slides (“Protocol for obtaining the empirical RTD...”)

Note: Time granularity based on μsec has been chosen to evaluate time duration of individual search iterations.

Note: Time measurement is based on Python `time.perf_counter_ns()`, which offers the largest available time resolution ($10\text{e-}9$ vs $10\text{e-}6$ for `time.process_time_ns()`). It must be noted the `perf_counter_ns()` also includes IDLE times, but especially when measuring the uf20-* instances with WalkSAT the gain in time resolution outweighs the potential missing restriction to process time only. Larger SAT instances should be measured using `time.process_time_ns()`.
[1]

Note: Only Section 3 Run-Time Distribution Evaluation expands on problem setup and restart time.

1 GWSAT

This section analyses the impact of the GWSAT parameters (restarts, iterations, wp) on the algorithm's operations against a selected set of CNF instances (uf20-01, uf20-02, uf50-01). The number of executions is fixed to 30. All statistics will be collected with a random seed defined for experiment repeatability.

The default behaviour of GSAT is to consider every variable in every iteration, calculate the effect of each of the potential flip and select the variable that provides the largest net_gain – which can also be negative. This greedy operation can lead to long plateaus and the algorithm getting stuck in a local minimum altogether. The default behaviour for the algorithm to escape from a local minimum is “restarting” after a fixed number of iterations. Such a restart creates a new variable assignment and the GSAT algorithm starts again. Considering that a classical GSAT step takes every variable into consideration in every iteration one might not speak of a solution intensification, despite GSAT's main metric of selecting the variable with the minimal net_gain.

By introducing the random walk parameter “wp” GWSAT provides a second option for the algorithm to escape a local minimum. “wp” represents the probability with which GWSAT will randomly select a variable currently participating in at least one unsatisfied clause and flip it, no matter how good and bad this move is. In iterations with a random walk, GWSAT will not evaluate the net_gain of any other variable. The random walk provides a means of diversifying the search space.

Therefore, a careful selection of the walk parameter “wp” can assist the GWSAT algorithm with:

1. Escape local minima
2. Reduce algorithm runtime

The other 2 parameters “restart” and “max_iterations” have a strong relationship. The parameter “max_iterations” defines how many search steps are executed before the algorithm starts over. While reducing “max_iterations” can upper bound the CPU cycles or time spent in a single search round, such change will result in an increase of restarts the GWSAT algorithm will have to go through. Therefore, depending on the size of the problem instance the parameters “restart” and “max_iterations” need to be carefully chosen, i.e. reducing “max_iterations” to less than the average iterations required to solve a given instance has to be avoided.

A large number of experiments will be executed and documented on the next pages to demonstrate the parameters' impact on the algorithm's operation. The parameters' impact will first be measured against algorithm search performance of the uf50-01 instance. Afterwards, found hypothesis will be validated against the small uf20 instances as well.

1.1 Evaluation of “wp”

Experiment	Instance	Executions	Restarts	Iterations	wp
1.1.0	uf50-01	30	10	1000	0.4
1.1.1	uf50-01	1	10	1000	0.001
1.1.2	uf50-01	1	10	1000	0.05
1.1.3	uf50-01	1	10	1000	0.2
1.1.4	uf50-01	1	10	1000	0.4
1.1.5	uf50-01	1	10	1000	0.6
1.1.6	uf50-01	1	10	1000	0.8
1.1.7	uf50-01	30	10	1000	0.001
1.1.8	uf50-01	30	10	1000	0.05
1.1.9	uf50-01	30	10	1000	0.2
1.1.10	uf50-01	30	10	1000	0.6
1.1.11	uf50-01	30	10	1000	0.8

Experiment 1.1.0 represents a baseline experiment against instance uf50-01 with parameters provided in the assignment brief. From the statistics we can see that on average 370 iterations are needed for search runs (restarts) where a solution is found. We also see that the longest search took 3379 iterations, that is 3 restarts after 1000 iterations without finding a solution plus another 379 iterations to solve the instance.

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
1.1.0	Total : 17718895 μ sec Mean : 590629 μ sec Success rate: 100.00%	Sum : 16 Min : 0 Max : 3 Mean : 0.5 Median: 0.0 STDEV : 0.8	Sum : 27072 Min : 23 Max : 3379 Mean : 902.4 Median: 548.5 STDEV : 821.9 VarCo : 0.911	Min : 23 Max : 876 Mean : 369.1 Median: 363.5 STDEV : 204.2 VarCo : 0.553

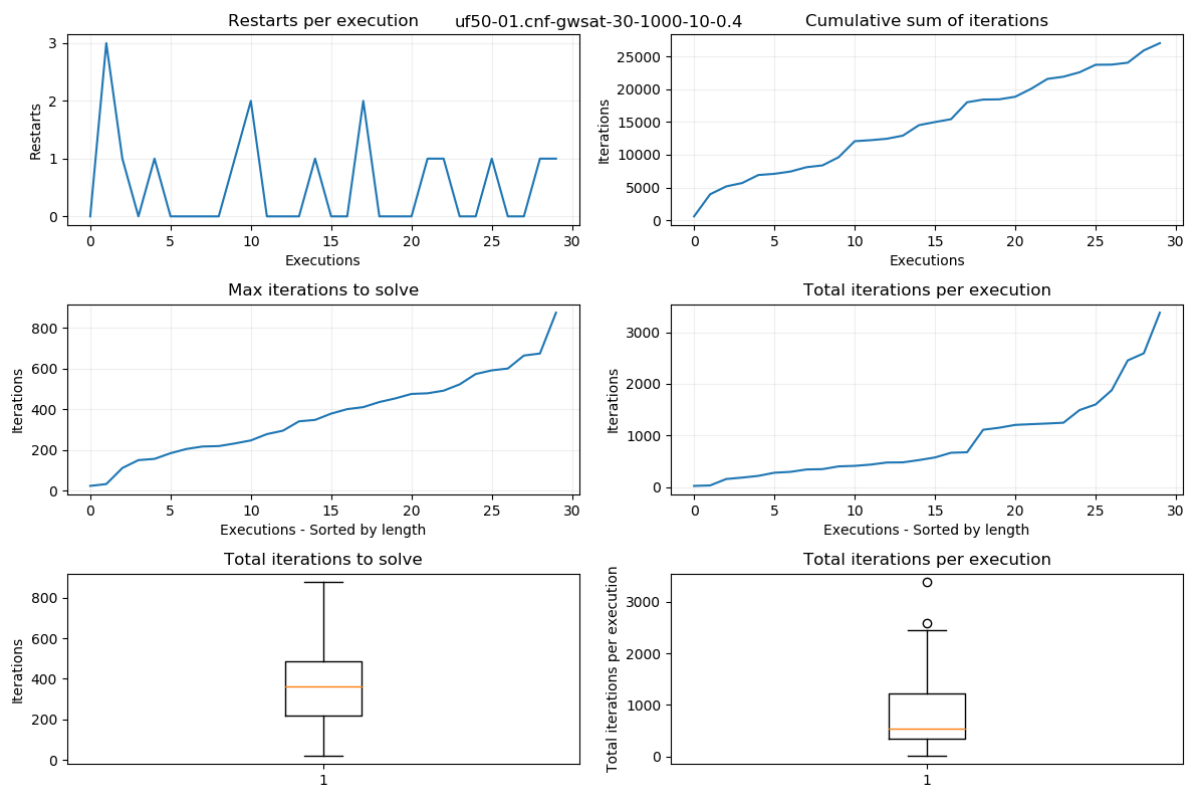
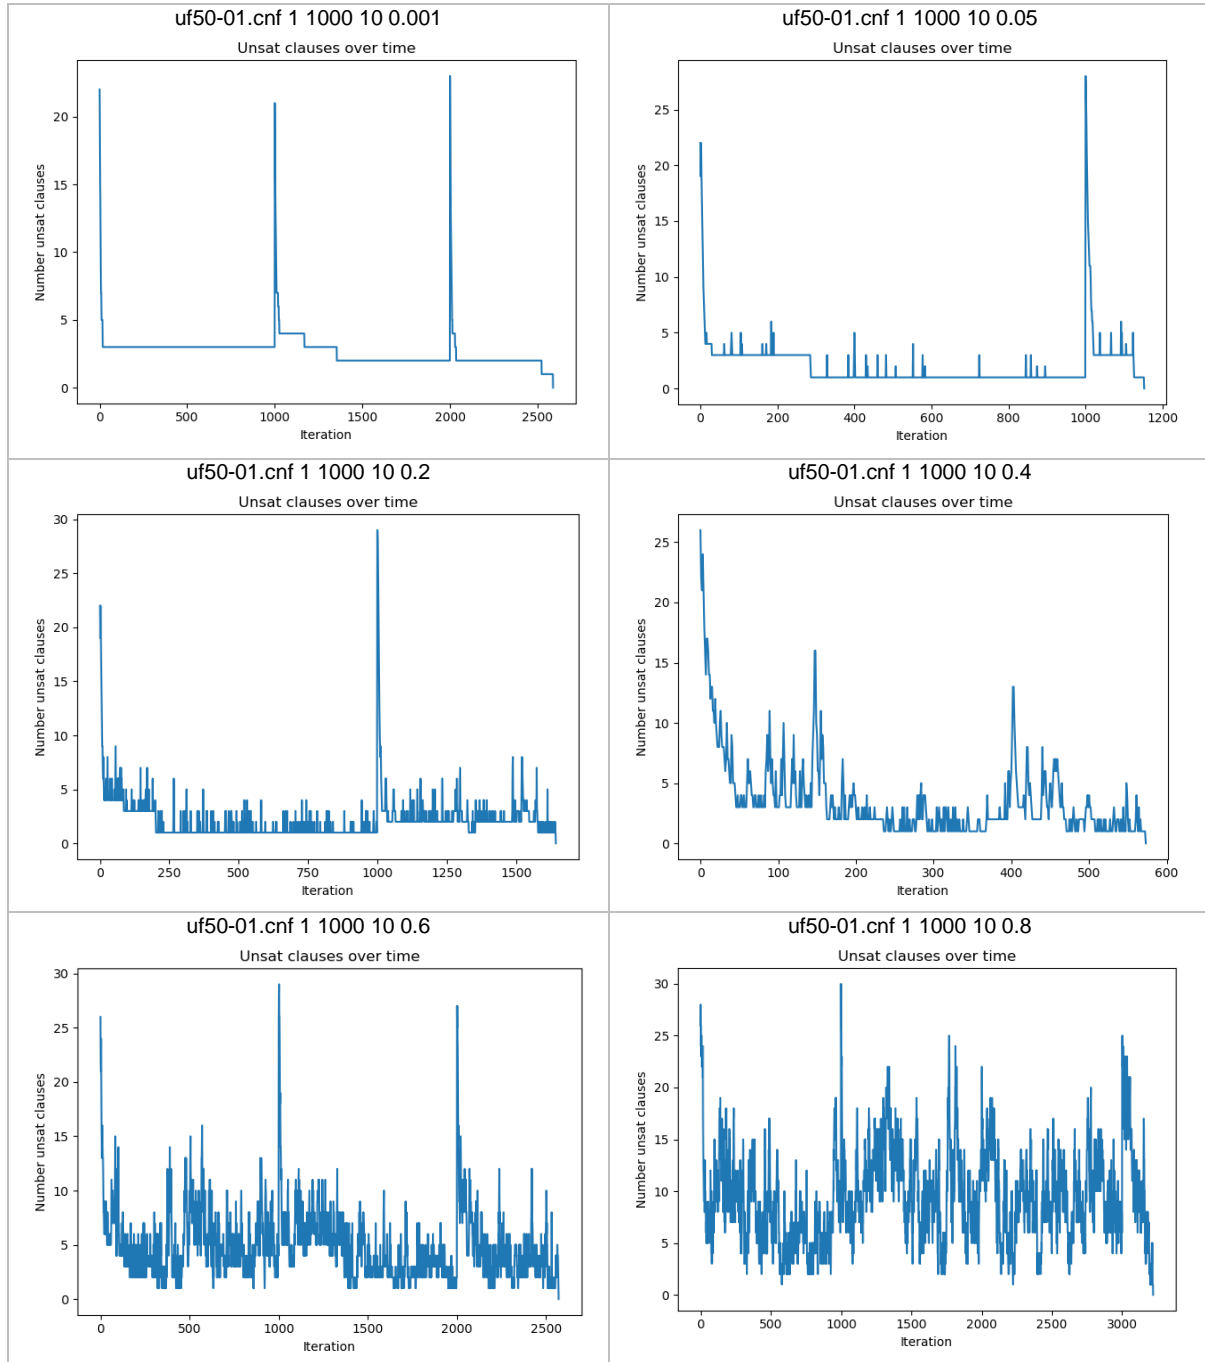


Figure 1 Restart and Iterations statistics for experiment 1.1.0

Next, experiments 1.1.1 to 1.1.6 will be run with a single execution in order to visualize the impact of parameter “wp” on the number of unsatisfied clauses during the search operation. I will demonstrate the impact of $wp = [0.001, 0.05, 0.2, 0.4, 0.6, 0.8]$. Given the default “max_iterations” of 1000 these “wp” probabilities will lead 0.01%, 5%, 20%, 40%, 60% and 80% expected random walk steps, or 1 out of 1000 to 800 out of 1000.



From experiment 1.1.1 ($wp = 0.001$) we can see that GWSAT is almost exclusively relying on a search restart to escape a local minimum. The same seems to be true for experiment 1.1.2. One notices a couple of spikes attributed to the random walk component of GWSAT, but with low “wp” probabilities the next iteration relies on classical GSAT operations which immediately reverts the random step and the algorithm is stuck in the local minimum again.

Experiments 1.1.4 seems to visualize a sweet spot for the random walk component in GWSAT. The random walk probability is high enough to ensure a low number of subsequent random walks to escape

a local minimum and avoid the immediate reversal of the random walk step. Experiment 1.1.3 is hard to judge. One notices few isolated subsequent steps resulting in a higher number of unsatisfied clauses, but in general the algorithm quickly drops back into the local minimum it tried to escape.

Experiments 1.1.5 and 1.1.6 more and more develop into an uncontrolled and purely random search and seem to be driven by less suitable “wp” values (0.6 and 0.8).

For comparison lets run experiments 1.1.7 to 1.1.11 with 30 executions. Please note that a 30 executions experiment with wp=0.4 was already executed as part of experiment 1.1.0. The results for experiment 1.1.0 are repeated for comparability.

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
1.1.0	Total : 17718895 μ sec	Sum : 16	Sum : 27072	
	Mean : 590629 μ sec	Min : 0	Min : 23	Min : 23
	Success rate: 100.00%	Max : 3	Max : 3379	Max : 876
		Mean : 0.5	Mean : 902.4	Mean : 369.1
	Total iter: 27072	Median: 0.0	Median: 548.5	Median: 363.5
	Avg iter time: 655 μ sec	STDEV : 0.8	STDEV : 821.9	STDEV : 204.2
			VarCo : 0.911	VarCo : 0.553
1.1.7	Total : 108571670 μ sec	Sum : 87	Sum : 95519	
	Mean : 3619055 μ sec	Min : 0	Min : 39	Min : 23
	Success rate: 96.67%	Max : 8	Max : 8860	Max : 902
		Mean : 3	Mean : 3293.8	Mean : 293.8
	Total iter: 95519	Median: 2	Median: 2589	Median: 242
	Avg iter time: 1027 μ sec	STDEV : 2.5	STDEV : 2580.1	STDEV : 248.9
			VarCo : 0.783	VarCo : 0.847
1.1.8	Total : 74070084 μ sec	Sum : 55	Sum : 65571	
	Mean : 2469002 μ sec	Min : 0	Min : 42	Min : 18
	Success rate: 96.67%	Max : 6	Max : 6740	Max : 961
		Mean : 1.9	Mean : 2261.1	Mean : 364.5
	Total iter: 65571	Median: 1	Median: 1790	Median: 193
	Avg iter time: 981 μ sec	STDEV : 1.9	STDEV : 1925.9	STDEV : 339.7
			VarCo : 0.852	VarCo : 0.932
1.1.9	Total : 39297546 μ sec	Sum : 35	Sum : 46080	
	Mean : 1309918 μ sec	Min : 0	Min : 20	Min : 20
	Success rate: 100.00%	Max : 6	Max : 6961	Max : 961
		Mean : 1.2	Mean : 1536	Mean : 369.3
	Total iter: 46080	Median: 1.0	Median: 1062.5	Median: 323.0
	Avg iter time: 853 μ sec	STDEV : 1.8	STDEV : 1875.4	STDEV : 255.7
			VarCo : 1.221	VarCo : 0.692
1.1.10	Total : 16773580 μ sec	Sum : 22	Sum : 37762	
	Mean : 559119 μ sec	Min : 0	Min : 120	Min : 120
	Success rate: 100.00%	Max : 6	Max : 6864	Max : 946
		Mean : 0.7	Mean : 1258.7	Mean : 525.4
	Total iter: 37762	Median: 0.0	Median: 677.5	Median: 571.0
	Avg iter time: 444 μ sec	STDEV : 1.3	STDEV : 1398.6	STDEV : 215.0
			VarCo : 1.111	VarCo : 0.409
1.1.11	Total : 33371557 μ sec	Sum : 85	Sum : 99583	
	Mean : 1112385 μ sec	Min : 0	Min : 328	Min : 52
	Success rate: 86.67%	Max : 9	Max : 9838	Max : 1000
		Mean : 3.3	Mean : 3830.1	Mean : 560.9
	Total iter: 99583	Median: 3.0	Median: 3218.5	Median: 562.5
	Avg iter time: 240 μ sec	STDEV : 2.7	STDEV : 2736.2	STDEV : 271.5
			VarCo : 0.714	VarCo : 0.484

By comparing absolute search times experiments 1.1.0 and 1.1.10 seem to be on par, but the total number of iterations required is substantially smaller in experiment 1.1.0. However, in experiment 1.1.10 a 60% of the iterations on average were driven by the random walk operation so that the number of classical GSAT iterations is almost identical between experiments 1.1.0 and 1.1.10. One should still note that the mean iterations to solve the problem is much smaller for experiment 1.1.0.

As noted above high probabilities for the random walk parameter “wp” lead to a diversification of the search space resulting in more states to be analysed.

Figure 2 visualizes the key performance indicators discussed above.




Figure 2 Total Iterations and Runtime in sec per wp parameter

Figure 3 visualizes how the selection of "wp" affects the runtime of a single search iteration. The larger the walk probability "wp" is the more iterations the algorithm will shortcut its classical GSAT protocol. Even for small problem instances with 20 variables not calculating the net_gain for each variable in every iteration shows a huge performance improvement.




Figure 3 Average Iteration time in μsec per wp parameter

This also highlights GWSATs biggest weakness: With moderate settings of "wp" its time complexity is close to $O(m \cdot n)$, where m is the number of search iterations and n the number of variables in a SAT instance.

1.2 Evaluation of “restart” and “max_iterations”

As already noted above the “restart” parameter is strongly affected by the “max_iterations” parameter. A search that cannot find a solution in a given run, will result in a search restart or terminate without result.

We have seen in section 1.1 that for experiment 1.1.0 the average number of iterations when solving instance uf50-01 is 369 and from Figure 1 one can read that in 27 out of 30 executions the search run finding a solution required 600 or less iterations. This provides already a framing for the next experiments where the number of iterations will be set to 200, 400, 600 and 800 respectively.

Experiment	Instance	Executions	Restarts	Iterations	wp
1.2.1	uf50-01	30	10	200	0.4
1.2.2	uf50-01	30	10	400	0.4
1.2.3	uf50-01	30	10	500	0.4
1.2.4	uf50-01	30	10	600	0.4
1.2.5	uf50-01	30	10	800	0.4
1.2.6	uf50-01	30	10	1200	0.4

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
1.2.1	Total : 18373665 µsec Mean : 612455 µsec Success rate: 83.33% Total iter: 18549 Avg iter time: 644 µsec	Sum : 78 Min : 0 Max : 9 Mean : 3.1 Median: 2 STDEV : 2.9	Sum : 18549 Min : 56 Max : 1892 Mean : 742.0 Median: 575 STDEV : 568.6 VarCo : 0.766	Min : 16 Max : 199 Mean : 118.0 Median: 115 STDEV : 52.6 VarCo : 0.446
1.2.2	Total : 17693343 µsec Mean : 589778 µsec Success rate: 100.00% Total iter: 27659 Avg iter time: 640 µsec	Sum : 55 Min : 0 Max : 9 Mean : 1.8 Median: 1.0 STDEV : 2.5	Sum : 27659 Min : 38 Max : 3929 Mean : 922.0 Median: 516.0 STDEV : 1035.5 VarCo : 1.123	Min : 38 Max : 391 Mean : 188.6 Median: 156.0 STDEV : 106.8 VarCo : 0.566
1.2.3	Total : 17732704 µsec Mean : 591090 µsec Success rate: 100.00% Total iter: 27107 Avg iter time: 654 µsec	Sum : 40 Min : 0 Max : 6 Mean : 1.3 Median: 1.0 STDEV : 1.6	Sum : 27107 Min : 109 Max : 3244 Mean : 903.6 Median: 717.0 STDEV : 789.2 VarCo : 0.873	Min : 70 Max : 477 Mean : 236.9 Median: 229.5 STDEV : 119.3 VarCo : 0.504
1.2.4	Total : 18565400 µsec Mean : 618846 µsec Success rate: 100.00% Total iter: 28772 Avg iter time: 645 µsec	Sum : 34 Min : 0 Max : 4 Mean : 1.1 Median: 1.0 STDEV : 1.3	Sum : 30573 Min : 57 Max : 3765 Mean : 1019.1 Median: 810.5 STDEV : 860.5 VarCo : 0.844	Min : 37 Max : 599 Mean : 279.1 Median: 251.5 STDEV : 163.3 VarCo : 0.585
1.2.5	Total : 19678288 µsec Mean : 655942 µsec Success rate: 100.00% Total iter: 30573 Avg iter time: 644 µsec	Sum : 27 Min : 0 Max : 4 Mean : 0.9 Median: 0.5 STDEV : 1.1	Sum : 27072 Min : 23 Max : 3379 Mean : 902.4 Median: 548.5 STDEV : 821.9 VarCo : 0.911	Min : 57 Max : 749 Mean : 299.1 Median: 201.0 STDEV : 224.2 VarCo : 0.749
1.2.6	Total : 18795113 µsec Mean : 626503 µsec Success rate: 100.00% Total iter: 29245 Avg iter time: 643 µsec	Sum : 11 Min : 0 Max : 3 Mean : 0.4 Median: 0.0 STDEV : 0.8	Sum : 29245 Min : 108 Max : 4721 Mean : 974.8 Median: 594.5 STDEV : 1078.0 VarCo : 1.106	Min : 63 Max : 1177 Mean : 534.8 Median: 535.0 STDEV : 343.5 VarCo : 0.642

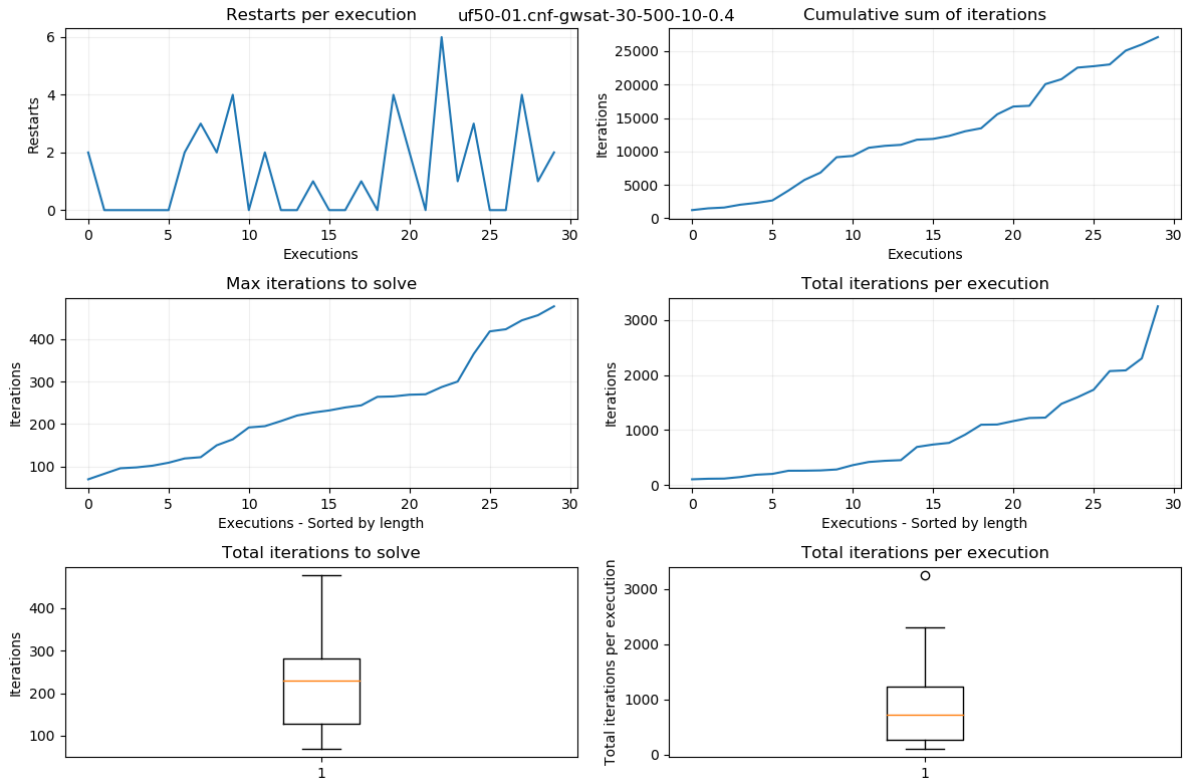


Figure 4 Restart and Iterations statistics for experiment 1.2.3

It seems that for solving instance `uf50-01` the limitation of “`max_iterations`” to 500 is a reasonable choice. The number of overall iterations to solve all 30 executions is only minimally larger than experiment 1.1.0, but the reduced “`max_iterations`” parameter may help in estimating worst case runtimes. Also with maximum 6 restarts per execution, this setup is still reasonably far away from terminating the search without finding a solution. With “`wp`” being fixed to 0.4 the computation time per iteration is also fixed.

It should be noted that larger SAT problems may require larger “`max_iterations`” and “`restart`” parameters.

1.3 Evaluating instances uf20-01 and uf20-02

After having analysed the impact of the parameters to the larger uf50-01 instance, a few validation checks are run against instances uf20-01 and uf20-02.

Experiment	Instance	Executions	Restarts	Iterations	wp
1.3.1	uf20-01	30	10	1000	0.2
1.3.2	uf20-01	30	10	1000	0.4
1.3.3	uf20-01	30	10	1000	0.6
1.3.4	uf20-02	30	10	1000	0.2
1.3.5	uf20-02	30	10	1000	0.4
1.3.6	uf20-02	30	10	1000	0.6

Experiment	Execution Overview	Overview	Restarts	Total Iterations	Iterations to solve
1.3.1	Total : 517798 μ sec Mean : 17259 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 2187 Min : 6 Max : 656 Mean : 72.9 Median: 31.0 STDEV : 120.2 VarCo : 1.649	Min : 6 Max : 656 Mean : 72.9 Median: 31.0 STDEV : 120.2 VarCo : 1.649
1.3.2	Total : 202389 μ sec Mean : 6746 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 1123 Min : 6 Max : 109 Mean : 37.4 Median: 27.0 STDEV : 31.1 VarCo : 0.83	Min : 6 Max : 109 Mean : 37.4 Median: 27.0 STDEV : 31.1 VarCo : 0.83
1.3.3	Total : 264895 μ sec Mean : 8829 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 1943 Min : 8 Max : 246 Mean : 64.8 Median: 47.0 STDEV : 51.7 VarCo : 0.799	Min : 8 Max : 246 Mean : 64.8 Median: 47.0 STDEV : 51.7 VarCo : 0.799
1.3.4	Total : 201304 μ sec Mean : 6710 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 781 Min : 4 Max : 118 Mean : 26.0 Median: 17.0 STDEV : 23.8 VarCo : 0.915	Min : 4 Max : 118 Mean : 26.0 Median: 17.0 STDEV : 23.8 VarCo : 0.915
1.3.5	Total : 198928 μ sec Mean : 6630 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 1060 Min : 9 Max : 124 Mean : 35.3 Median: 24.5 STDEV : 29.0 VarCo : 0.82	Min : 9 Max : 124 Mean : 35.3 Median: 24.5 STDEV : 29.0 VarCo : 0.82
1.3.6	Total : 153340 μ sec Mean : 5111 μ sec Success rate: 100.00%		Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 1155 Min : 4 Max : 121 Mean : 38.5 Median: 25.0 STDEV : 29.7 VarCo : 0.771	Min : 4 Max : 121 Mean : 38.5 Median: 25.0 STDEV : 29.7 VarCo : 0.771

Using a random walk probability of 0.4 against the smaller SAT problem instances generally proves useful as well. The “max_iterations” parameter could also be safely reduced to 200 without affecting any of the above results. With small problem instances also the total search execution but also the runtime per iteration is proportionally lower.

1.4 Conclusion

The previous experiments have given insights into how the operation of GWSAT is affected by its set of parameters “restart”, “max_iterations” and “wp”. A too low setting for “wp” leads to the random walk component being invisible or reverted immediately by a following classical GSAT step. It seems that with a “wp” value around 0.4 the probability of executing multiple random walk steps in sequence is large enough to be able to escape local minima.

The “wp” probabilities between 0.4 and 0.6 offer an interesting trade-off between a higher number of total iterations, but also more iterations being short-cut by a random walk. “wp” values higher than 0.6 seem to end in a close(r) to random search.

Selecting an appropriate value for “max_iterations” depends on the problem size to work with and how many iterations are usually required to solve a problem. If pre-maturely a search is stopped in favour of a restart, the search needs to work down the number of unsatisfied clauses after re-initializing and consuming extra time.

As noted above, GWSAT’s tendency to evaluate the flip of every variable for classical GSAT steps, makes it scale very badly to large problem instances.

1.5 Debug GWSAT

To debug the operation of the GWSAT implementation set the DEBUG flag to TRUE in file ‘Leske_183658_GWSAT.py’.

This will create an output as follows:

```
...
total_unsat:      2
candidate_vars: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20]
gains:           [-5, -1, -3, 0, -1, -3, -3, -3, -2, -2, -1, -1, 0, -2, -1, -1, -1, -4, -2, -4]
flip_var:        4

total_unsat:      2
candidate_vars: [10]
gains:           [-2]
flip_var:        10

total_unsat:      4
candidate_vars: [18]
gains:           [-3]
flip_var:        18

total_unsat:      7
candidate_vars: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20]
gains:           [-4, 0, -3, 0, 0, -1, -3, 0, 1, 1, 0, -2, 0, 0, -1, 0, -1, 3, -1, -4]
flip_var:        18
...
```

This snippet shows 4 iterations. In iteration 1 a classical GSAT step was executed. All variables acted as candidates and for each a flip was simulated and the net_gain calculated. Variables 4 and 13 had the best net_gain of 0. A random tie breaker selected variable 4 to be flipped.

Next we see that the total_unsat counter remained 2 due to 0 net_gain of the flipped variable. This iteration represents a random walk iteration as the candidate list only contains a single variable.

Iteration 3 represents a random walk step as well. The net_gain of the random operation is -3. The classical GSAT step in the 4th iteration shown above finds that the best net_gain is flipping back the previously randomly selected variable 18.

2 WalkSAT/SKC

In contrast to GSAT the original WalkSAT algorithm according to SKC differs in two essential ways:

1. In each iteration an unsatisfied clause is randomly selected. Only variables included in this clause will be flip candidates for the given iteration.
2. The metric to decide which variable to flip is based on the “negative gain” counter, which represents the number of clauses that were satisfied before the intended flip and would be unsatisfied thereafter.

While the focus on a single unsatisfied clause only in any given iteration significantly reduces computing requirements as the impact of flipping less variables has to be analysed, WalkSAT/SKC might miss making the highest net_gain flip possible (as GWSAT does). With its focus on “negative gain” as variable selection metric, WalkSAT/SKC ensures making the least destructive flip in any given iteration. If one or more variables result in a flip with a zero “negative gain”, WalkSAT randomly picks one of these variables to ensure the number of unsatisfied clauses does not increase. If no such variable exists WalkSAT/SKC selects with a probability “p” a random variable from the selected unsatisfied clause, or with probability 1-p the variable that will result in the lowest “negative_gain”. Overall, this extremely careful variable selection metric helps WalkSAT/SKC to compensate for its much smaller search space per iteration. The variable selection process of WalkSAT/SKC can be described to intensifying the search space due to its focus on unsatisfied clauses only. The random operation according to parameter “p” introduce randomness into the search process.

By leveraging a Tabu list the search process can track the history of variable flips and essentially block single variables to be flipped within in a specified number of iterations “tl”. This also helps the algorithm to escape local minima by avoiding flipping back and forth the same variable once such local minimum has been reached. Blocking recently flipped variables from being selected again within a given iteration interval also increases the diversification of the search process.

A large number of experiments will be executed and documented on the next pages to demonstrate the parameters’ impact on the algorithm’s operation. The parameters’ impact will first be measured against algorithm search performance of the uf50-01 instance. Afterwards, found hypothesis will be validated against the small uf20 instances as well.

2.1 Evaluation of “p”

Experiment	Instance	Executions	Restarts	Iterations	wp	tl
2.1.0	uf50-01	30	10	1000	0.4	5
2.1.1	uf50-01	1	10	1000	0.001	5
2.1.2	uf50-01	1	10	1000	0.05	5
2.1.3	uf50-01	1	10	1000	0.2	5
2.1.4	uf50-01	1	10	1000	0.4	5
2.1.5	uf50-01	1	10	1000	0.6	5
2.1.6	uf50-01	1	10	1000	0.8	5
2.1.7	uf50-01	30	10	1000	0.001	5
2.1.8	uf50-01	30	10	1000	0.05	5
2.1.9	uf50-01	30	10	1000	0.2	5
2.1.10	uf50-01	30	10	1000	0.6	5
2.1.11	uf50-01	30	10	1000	0.8	5

Experiment 2.1.0 represents a baseline experiment against instance uf50-01 with parameters provided in the assignment brief. From the statistics we can see that on average 300 iterations are needed in search runs that find a solution. We also see that the longest search took 1497 iterations, that is 1 restart after 1000 iterations without finding a solution plus another 497 iterations to solve the instance. It is also immediately noticeable that the focus on a single unsatisfied clause lets the algorithm select a variable in much less time than a classical GSAT step.

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
2.1.0	Total : 834222 μ sec Mean : 27807 μ sec Success rate: 100.00%	Sum : 4 Min : 0 Max : 1 Mean : 0.1 Median: 0.0 STDEV : 0.3	Sum : 12990 Min : 15 Max : 1497 Mean : 433 Median: 283.5 STDEV : 419.5 VarCo : 0.969	Min : 15 Max : 853 Mean : 299.7 Median: 266.5 STDEV : 222.7 VarCo : 0.743

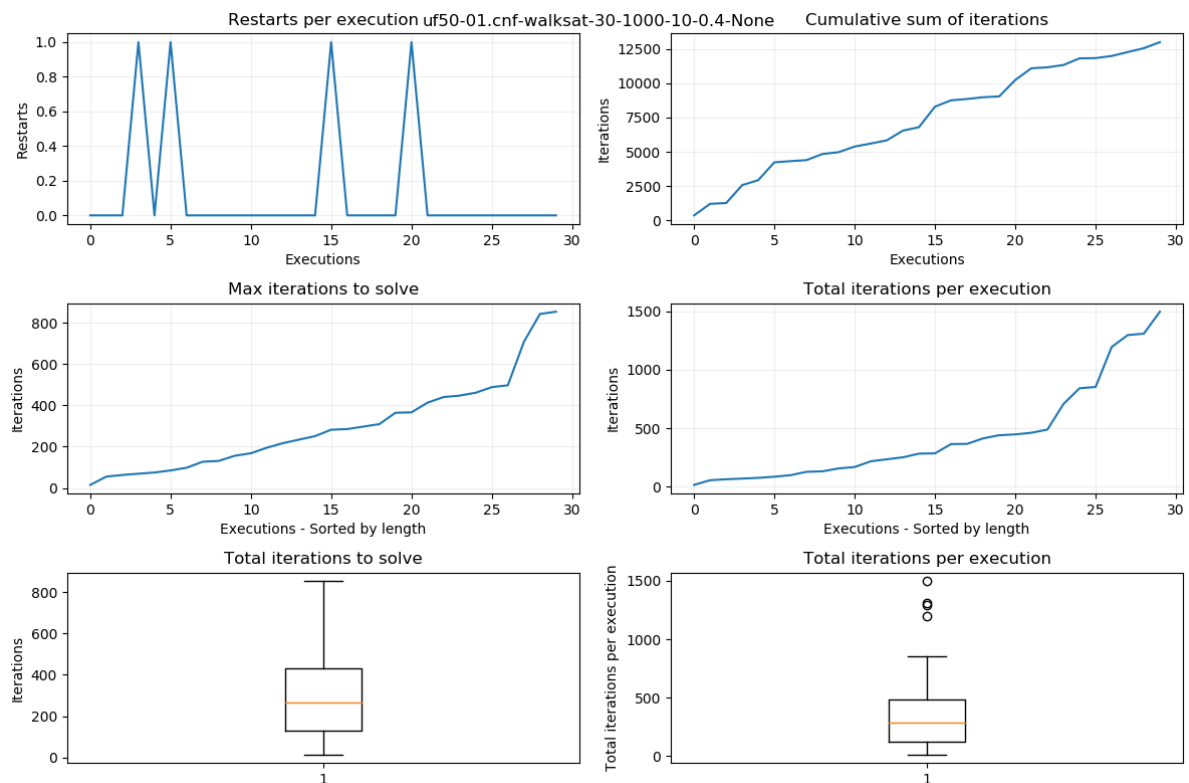
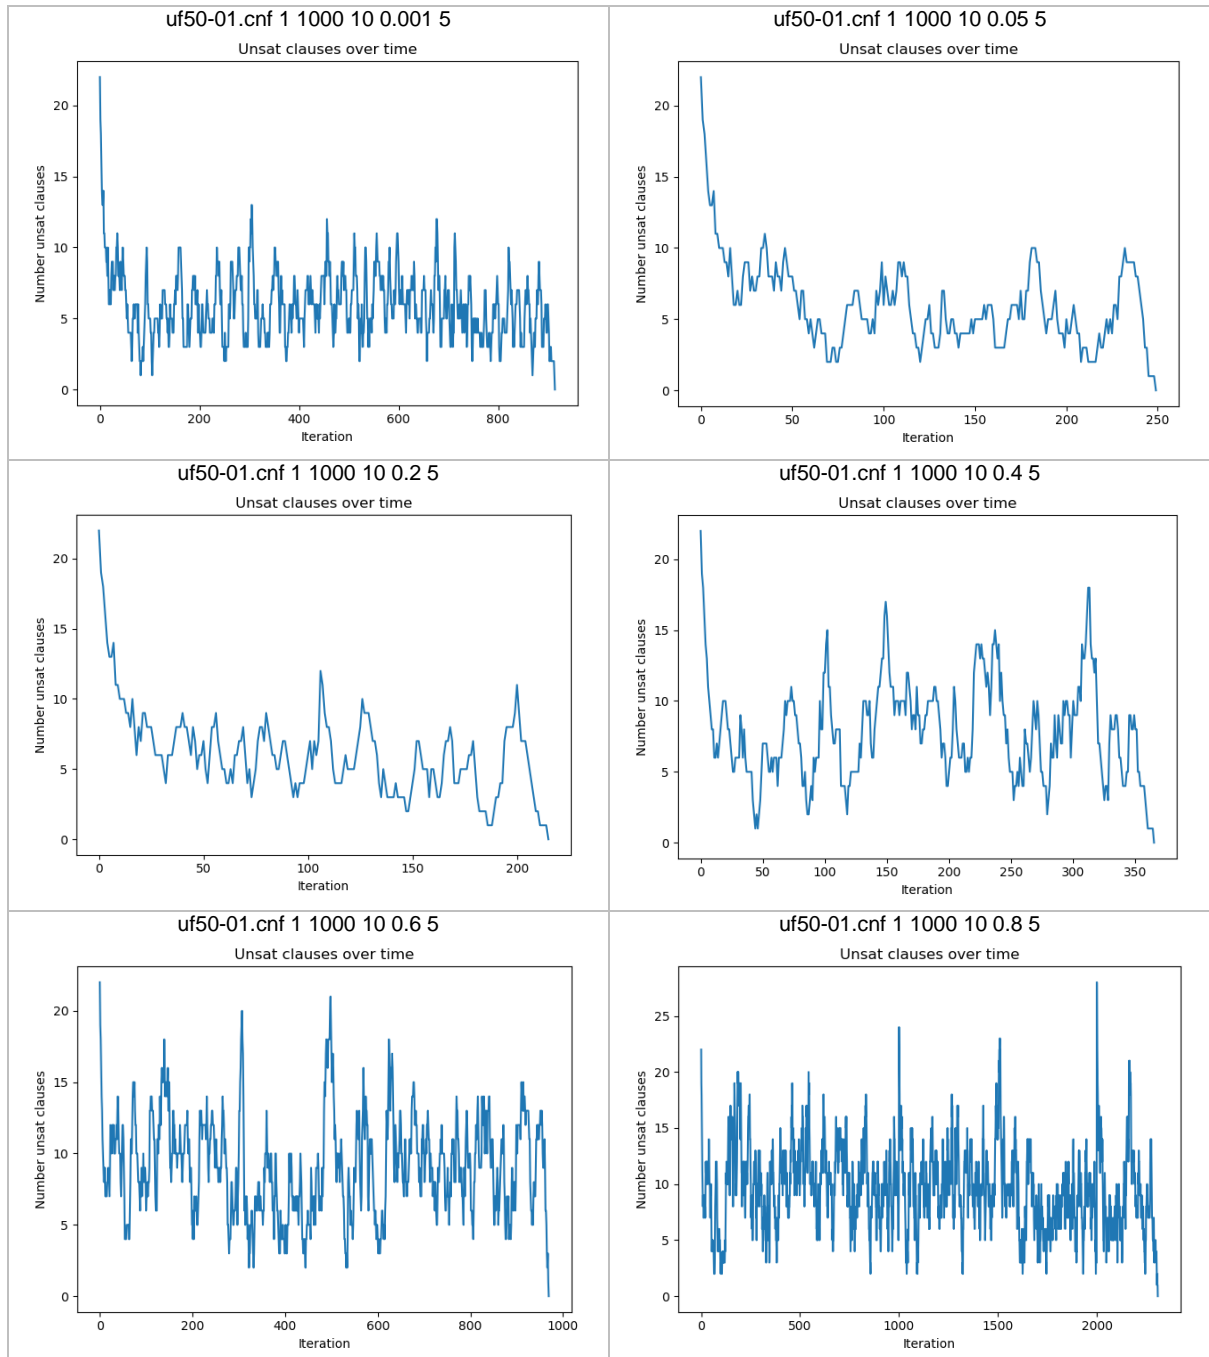


Figure 5 Restart and Iterations statistics for experiment 2.1.0

All statistics as well as the overall search time to complete 30 executions show significant improvement over the GWSAT search

As done for GWSAT, experiments 2.1.1 to 2.1.6 will be run with a single execution in order to visualize the impact of parameter “p” on the number of unsatisfied clauses during the search operation. The same values as for GWSAT “wp” will be used for WalkSAT’s “p” parameter:
[0.001, 0.05, 0.2, 0.4, 0.6, 0.8].



Considering that WalkSAT/SK solves instance uf50-01 with “p” = 0 in less than 1000 iterations as well, the “p” setting of 0.001 is expected to have extremely limited influence only – if at all.

Despite the higher number of iterations in experiment 2.1.1 and the squeezed graph representation, the actual graph is very close to the one of experiment 2.1.2 (p=0.05). Experiments 2.1.5 and 2.1.6

demonstrate a similar influence of a very high “wp” value for GSAT. The careful selection of variables according to the minimum “negative_gain” gets mostly replaced by random steps.

For comparison lets run experiments 2.1.7 to 2.1.11 with 30 executions and see a general trend when running certain configurations for multiple executions. Please note that a 30 executions experiment with p=0.4 was already executed as part of experiment 2.1.0. The results for experiment 2.1.0 are repeated for comparability.

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
2.1.0	Total : 834222 μ sec	Sum : 4	Sum : 12990	
	Mean : 27807 μ sec	Min : 0	Min : 15	Min : 15
	Success rate: 100.00%	Max : 1	Max : 1497	Max : 853
		Mean : 0.1	Mean : 433	Mean : 299.7
	Total iter: 12990	Median: 0.0	Median: 283.5	Median: 266.5
2.1.7	Avg iter time: 64 μ sec	STDEV : 0.3	STDEV : 419.5	STDEV : 222.7
			VarCo : 0.969	VarCo : 0.743
	Total : 610617 μ sec	Sum : 0	Sum : 9463	
	Mean : 20353 μ sec	Min : 0	Min : 30	Min : 30
	Success rate: 100.00%	Max : 0	Max : 980	Max : 980
2.1.8		Mean : 0	Mean : 315.4	Mean : 315.4
	Total iter: 9463	Median: 0.0	Median: 218.5	Median: 218.5
	Avg iter time: 65 μ sec	STDEV : 0.0	STDEV : 263.2	STDEV : 263.2
			VarCo : 0.835	VarCo : 0.835
	Total : 606179 μ sec	Sum : 1	Sum : 9231	
2.1.9	Mean : 20205 μ sec	Min : 0	Min : 66	Min : 66
	Success rate: 100.00%	Max : 1	Max : 1128	Max : 747
		Mean : 0.0	Mean : 307.7	Mean : 274.4
	Total iter: 9231	Median: 0.0	Median: 271.0	Median: 259.0
	Avg iter time: 66 μ sec	STDEV : 0.2	STDEV : 221.1	STDEV : 160.2
2.1.10			VarCo : 0.719	VarCo : 0.584
	Total : 826685 μ sec	Sum : 3	Sum : 12395	
	Mean : 27556 μ sec	Min : 0	Min : 25	Min : 25
	Success rate: 100.00%	Max : 1	Max : 1293	Max : 806
		Mean : 0.1	Mean : 413.2	Mean : 313.2
2.1.11	Total iter: 12395	Median: 0.0	Median: 304.0	Median: 276.0
	Avg iter time: 67 μ sec	STDEV : 0.3	STDEV : 329.4	STDEV : 188.8
			VarCo : 0.797	VarCo : 0.603
	Total : 1692692 μ sec	Sum : 11	Sum : 26330	
	Mean : 56423 μ sec	Min : 0	Min : 38	Min : 38
2.1.12	Success rate: 100.00%	Max : 3	Max : 3495	Max : 970
		Mean : 0.4	Mean : 877.7	Mean : 511
	Total iter: 26330	Median: 0.0	Median: 738.0	Median: 543.5
	Avg iter time: 64 μ sec	STDEV : 0.7	STDEV : 699.1	STDEV : 295.3
			VarCo : 0.797	VarCo : 0.578
2.1.13	Total : 4040370 μ sec	Sum : 40	Sum : 52646	
	Mean : 134679 μ sec	Min : 0	Min : 68	Min : 68
	Success rate: 96.67%	Max : 4	Max : 4852	Max : 979
		Mean : 1.4	Mean : 1815.4	Mean : 436.1
	Total iter: 52646	Median: 1	Median: 1618	Median: 433
2.1.14	Avg iter time: 64 μ sec	STDEV : 1.2	STDEV : 1159.4	STDEV : 279.3
			VarCo : 0.639	VarCo : 0.641

Experiments 2.1.7 to 2.1.11 confirm that WalkSAT generally works best with very small “p” values for 2 reasons:

1. It is able to honour its main metric of “negative_gain” for variable selection.
2. The random walk based on “p” does not provide any noticeable iteration runtime improvements as in any case only the variables of a single unsatisfied clause are considered as flip candidate in any given iteration. In addition, this list of potential flip candidates may be further reduced if one or more variables are tabued in a given iteration. Search runtime per iteration is constant.

Experiments 2.1.9 to 2.1.11 (and 2.1.0) illustrate that the increase in total number of iterations immediately translates into longer search runtime (for the reason given above).

Figure 5 visualizes the key performance indicators discussed above.




Figure 6 Total Iterations and Runtime in sec per wp parameter

2.2 Evaluation of “restart” and “max_iterations”

The trade-off between the parameters “restart” and “max_iterations” have been discussed in section 1.2 already. The same concepts apply to WalkSAT:

1. A decrease in “max_iterations” will result in more restarts (given the parameter is reduced to less than the longest run seen for a given instance).
2. Each restart has a processing cost in creating a new solution assignment and restart the search.

A few selected experiments will be executed for documenting the impact of changing the “max_iterations” parameter for the current problem instance uf50-01.

Experiment	Instance	Executions	Restarts	Iterations	wp	tl
2.2.1	uf50-01	30	10	400	0.001	5
2.2.2	uf50-01	30	10	600	0.001	5
2.2.3	uf50-01	30	10	800	0.001	5
2.2.4	uf50-01	30	10	400	0.05	5
2.2.5	uf50-01	30	10	600	0.05	5
2.2.6	uf50-01	30	10	800	0.05	5

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
2.2.1	Total : 652134 μ sec Mean : 21737 μ sec Success rate: 100.00% Total iter: 10058 Avg iter time: 65 μ sec	Sum : 12 Min : 0 Max : 3 Mean : 0.4 Median: 0.0 STDEV : 0.7	Sum : 10058 Min : 34 Max : 1327 Mean : 335.3 Median: 194.0 STDEV : 310.6 VarCo : 0.926	Min : 34 Max : 357 Mean : 175.3 Median: 153.0 STDEV : 95.3 VarCo : 0.544
2.2.2	Total : 440489 μ sec Mean : 14682 μ sec Success rate: 100.00% Total iter: 6747 Avg iter time: 65 μ sec	Sum : 2 Min : 0 Max : 1 Mean : 0.1 Median: 0.0 STDEV : 0.3	Sum : 6747 Min : 28 Max : 882 Mean : 224.9 Median: 176.0 STDEV : 185.3 VarCo : 0.824	Min : 26 Max : 449 Mean : 184.9 Median: 170.5 STDEV : 117.1 VarCo : 0.633
2.2.3	Total : 560000 μ sec Mean : 18666 μ sec Success rate: 100.00% Total iter: 8386 Avg iter time: 67 μ sec	Sum : 2 Min : 0 Max : 1 Mean : 0.1 Median: 0.0 STDEV : 0.3	Sum : 8386 Min : 22 Max : 944 Mean : 279.5 Median: 192.5 STDEV : 246.6 VarCo : 0.882	Min : 22 Max : 671 Mean : 226.2 Median: 171.0 STDEV : 184.4 VarCo : 0.815
2.2.4	Total : 664938 μ sec Mean : 22164 μ sec Success rate: 100.00% Total iter: 10246 Avg iter time: 65 μ sec	Sum : 12 Min : 0 Max : 3 Mean : 0.4 Median: 0.0 STDEV : 0.7	Sum : 10246 Min : 44 Max : 1418 Mean : 341.5 Median: 263.5 STDEV : 296.6 VarCo : 0.868	Min : 23 Max : 375 Mean : 181.5 Median: 172.5 STDEV : 111.6 VarCo : 0.615
2.2.5	Total : 605944 μ sec Mean : 20198 μ sec Success rate: 100.00% Total iter: 9311 Avg iter time: 65 μ sec	Sum : 3 Min : 0 Max : 1 Mean : 0.1 Median: 0.0 STDEV : 0.3	Sum : 9311 Min : 49 Max : 779 Mean : 310.4 Median: 260.0 STDEV : 191.6 VarCo : 0.617	Min : 49 Max : 574 Mean : 250.4 Median: 246.0 STDEV : 134.1 VarCo : 0.536
2.2.6	Total : 579269 μ sec Mean : 19308 μ sec Success rate: 100.00% Total iter: 8830 Avg iter time: 66 μ sec	Sum : 3 Min : 0 Max : 2 Mean : 0.1 Median: 0.0 STDEV : 0.4	Sum : 8830 Min : 30 Max : 1839 Mean : 294.3 Median: 236.5 STDEV : 342.2 VarCo : 1.163	Min : 30 Max : 447 Mean : 214.3 Median: 230.5 STDEV : 120.1 VarCo : 0.56

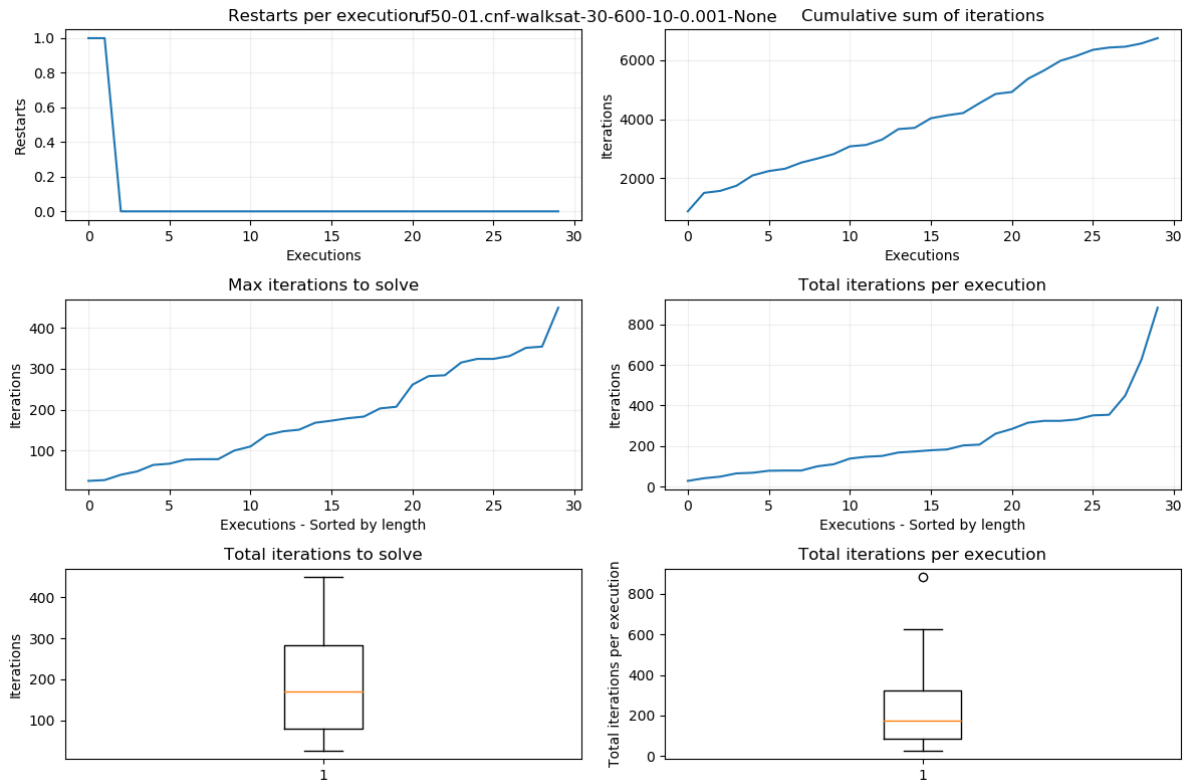


Figure 7 Restart and Iterations statistics for experiment 2.2.2

It seems that choosing a very small “p” value and reducing “max_iterations” to 600 worked extremely good in solving 30 searches for instance uf50-01.

2.3 Evaluation of “tl”

The parameter “tl” specifies the number of iterations a flipped variable is tabued until it can be flipped again. When used in conjunction with WalkSAT, the Tabu list effectively avoids that the algorithm is flipping the same variable back and forth once a local minimum has been reached and therefore ensures the search space’s diversification.

The parameter “tl” should be coupled to the number of variables in a given problem instance. Selecting “tl” too big, will result in too many variables being blocked and therefore limiting search progress. Selecting “tl” too small may lead the search into the same local minimum again.

For problem instance uf50-01 the “tl” setting of 5 achieved the best results. This represents 10% of the variables.

2.4 Evaluating instances uf20-01 and uf20-02

Lastly, experiments 2.4.1 to 2.4.8 measure the performance and conclusions from above against instances uf20-01 and uf20-02. The “restarts” and “max_iterations” parameters will be safely reduced to 3 and 200 respectively. “p” is set set 0.001. “tl” will be evaluated against a short set of values, mainly to see if the intuition of 10% variables blocked by the tabu list holds true. Experiments 2.4.1 and 2.4.5 represent baseline runs according to the assignment brief.

Experiment	Instance	Executions	Restarts	Iterations	wp	tl
2.4.1	uf20-01	30	10	1000	0.4	5
2.4.2	uf20-01	30	3	200	0.001	2
2.4.3	uf20-01	30	3	200	0.05	2
2.4.4	uf20-01	30	3	200	0.001	3
2.4.5	uf20-01	30	3	200	0.05	3
2.4.6	uf20-01	30	3	200	0.001	5
2.4.7	uf20-01	30	3	200	0.05	5

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
2.4.1	Total : 55916 μ sec	Sum : 0	Sum : 1086	
	Mean : 1863 μ sec	Min : 0	Min : 7	Min : 7
	Success rate: 100.00%	Max : 0	Max : 82	Max : 82
		Mean : 0	Mean : 36.2	Mean : 36.2
	Total iter: 1086	Median: 0.0	Median: 27.0	Median: 27.0
2.4.2	Avg iter time: 51 μ sec	STDEV : 0.0	STDEV : 23.2	STDEV : 23.2
			VarCo : 0.64	VarCo : 0.64
	Total : 54793 μ sec	Sum : 0	Sum : 980	
	Mean : 1826 μ sec	Min : 0	Min : 4	Min : 4
	Success rate: 100.00%	Max : 0	Max : 121	Max : 121
2.4.3		Mean : 0	Mean : 32.7	Mean : 32.7
	Total iter: 980	Median: 0.0	Median: 19.5	Median: 19.5
	Avg iter time: 56 μ sec	STDEV : 0.0	STDEV : 28.7	STDEV : 28.7
			VarCo : 0.878	VarCo : 0.878
	Total : 42507 μ sec	Sum : 0	Sum : 689	
2.4.4	Mean : 1416 μ sec	Min : 0	Min : 3	Min : 3
	Success rate: 100.00%	Max : 0	Max : 91	Max : 91
		Mean : 0	Mean : 23.0	Mean : 23.0
	Total iter: 689	Median: 0.0	Median: 17.0	Median: 17.0
	Avg iter time: 62 μ sec	STDEV : 0.0	STDEV : 22.2	STDEV : 22.2
2.4.5			VarCo : 0.967	VarCo : 0.967
	Total : 41519 μ sec	Sum : 0	Sum : 761	
	Mean : 1383 μ sec	Min : 0	Min : 3	Min : 3
	Success rate: 100.00%	Max : 0	Max : 85	Max : 85
		Mean : 0	Mean : 25.4	Mean : 25.4
2.4.6	Total iter: 761	Median: 0.0	Median: 17.5	Median: 17.5
	Avg iter time: 55 μ sec	STDEV : 0.0	STDEV : 22.9	STDEV : 22.9
			VarCo : 0.902	VarCo : 0.902
	Total : 58666 μ sec	Sum : 0	Sum : 1139	
	Mean : 1955 μ sec	Min : 0	Min : 8	Min : 8
2.4.7	Success rate: 100.00%	Max : 0	Max : 124	Max : 124
		Mean : 0	Mean : 38.0	Mean : 38.0
	Total iter: 1139	Median: 0.0	Median: 27.5	Median: 27.5
	Avg iter time: 52 μ sec	STDEV : 0.0	STDEV : 30.1	STDEV : 30.1
			VarCo : 0.792	VarCo : 0.792

2.4.6	Total : 45661 μ sec	Sum : 0	Sum : 911	
	Mean : 1522 μ sec	Min : 0	Min : 8	Min : 8
	Success rate: 100.00%	Max : 0	Max : 195	Max : 195
		Mean : 0	Mean : 30.4	Mean : 30.4
	Total iter: 911	Median: 0.0	Median: 16.5	Median: 16.5
	Avg iter time: 50 μ sec	STDEV : 0.0	STDEV : 36.8	STDEV : 36.8
			VarCo : 1.212	VarCo : 1.212
2.4.7	Total : 55869 μ sec	Sum : 0	Sum : 1139	
	Mean : 1862 μ sec	Min : 0	Min : 6	Min : 6
	Success rate: 100.00%	Max : 0	Max : 195	Max : 195
		Mean : 0	Mean : 38.0	Mean : 38.0
	Total iter: 1139	Median: 0.0	Median: 26.0	Median: 26.0
	Avg iter time: 49 μ sec	STDEV : 0.0	STDEV : 36.5	STDEV : 36.5
			VarCo : 0.961	VarCo : 0.961

The 3 most promising configurations will now be run with 1000 executions.

Experiment	Instance	Executions	Restarts	Iterations	wp	tl
2.4.8	uf20-01	1000	3	200	0.05	2
2.4.9	uf20-01	1000	3	200	0.001	3
2.4.10	uf20-01	1000	3	200	0.001	5

Experiment	Execution Overview	Restarts	Total Iterations	Iterations to solve
2.4.8	Total : 1728190 μ sec	Sum : 0	Sum : 31156	
	Mean : 1728 μ sec	Min : 0	Min : 1	Min : 1
	Success rate: 100.00%	Max : 0	Max : 247	Max : 189
		Mean : 0	Mean : 31.2	Mean : 30.8
	Total iter: 31156	Median: 0.0	Median: 23.0	Median: 23.0
	Avg iter time: 55 μ sec	STDEV : 0.0	STDEV : 26.9	STDEV : 25.4
			VarCo : 0.862	VarCo : 0.825
2.4.9	Total : 1683118 μ sec	Sum : 0	Sum : 30983	
	Mean : 1683 μ sec	Min : 0	Min : 2	Min : 2
	Success rate: 100.00%	Max : 0	Max : 176	Max : 176
		Mean : 0	Mean : 31.0	Mean : 31.0
	Total iter: 30983	Median: 0.0	Median: 24.0	Median: 24.0
	Avg iter time: 54 μ sec	STDEV : 0.0	STDEV : 25.9	STDEV : 25.9
			VarCo : 0.837	VarCo : 0.837
2.4.10	Total : 1674545 μ sec	Sum : 2	Sum : 34838	
	Mean : 1674 μ sec	Min : 0	Min : 3	Min : 3
	Success rate: 100.00%	Max : 1	Max : 223	Max : 195
		Mean : 0.0	Mean : 34.8	Mean : 34.4
	Total iter: 34838	Median: 0.0	Median: 26.0	Median: 26.0
	Avg iter time: 48 μ sec	STDEV : 0.0	STDEV : 30.8	STDEV : 29.7
			VarCo : 0.884	VarCo : 0.863

The above experiments confirm that also for very small instances like uf20-* a very low “p” value, e.g. 0.001 or 0.05, provides best results. The experiments 2.4.8 to 2.4.10 confirm that “tl” values of 2 or 3 (10%, or 15%) result in the least total iterations.

Experiment	Instance	Executions	Restarts	Iterations	wp	tl
2.4.11	uf20-02	30	10	1000	0.4	5
2.4.12	uf20-02	30	3	200	0.001	2
2.4.13	uf20-02	30	3	200	0.05	2
2.4.14	uf20-02	30	3	200	0.001	3
2.4.15	uf20-02	30	3	200	0.05	3
2.4.16	uf20-02	30	3	200	0.001	5
2.4.17	uf20-02	30	3	200	0.05	5
2.4.18	uf20-02	1000	3	200	0.001	2
2.4.19	uf20-02	1000	3	200	0.001	5


Experiment	Execution Overview Overview	Restarts	Total Iterations	Iterations to solve
2.4.11	Total : 50844 μ sec Mean : 1694 μ sec Success rate: 100.00% Total iter: 989 Avg iter time: 51 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 989 Min : 3 Max : 129 Mean : 33.0 Median: 30.5 STDEV : 24.9 VarCo : 0.755	Min : 3 Max : 129 Mean : 33.0 Median: 30.5 STDEV : 24.9 VarCo : 0.755
2.4.12	Total : 30523 μ sec Mean : 1017 μ sec Success rate: 100.00% Total iter: 455 Avg iter time: 67 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 455 Min : 5 Max : 34 Mean : 15.2 Median: 15.0 STDEV : 5.8 VarCo : 0.385	Min : 5 Max : 34 Mean : 15.2 Median: 15.0 STDEV : 5.8 VarCo : 0.385
2.4.13	Total : 30998 μ sec Mean : 1033 μ sec Success rate: 100.00% Total iter: 465 Avg iter time: 67 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 465 Min : 1 Max : 46 Mean : 15.5 Median: 13.5 STDEV : 9.1 VarCo : 0.587	Min : 1 Max : 46 Mean : 15.5 Median: 13.5 STDEV : 9.1 VarCo : 0.587
2.4.14	Total : 32556 μ sec Mean : 1085 μ sec Success rate: 100.00% Total iter: 487 Avg iter time: 67 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 487 Min : 3 Max : 34 Mean : 16.2 Median: 14.5 STDEV : 9.0 VarCo : 0.556	Min : 3 Max : 34 Mean : 16.2 Median: 14.5 STDEV : 9.0 VarCo : 0.556
2.4.15	Total : 32456 μ sec Mean : 1081 μ sec Success rate: 100.00% Total iter: 540 Avg iter time: 60 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 540 Min : 5 Max : 44 Mean : 18 Median: 16.5 STDEV : 10.5 VarCo : 0.584	Min : 5 Max : 44 Mean : 18 Median: 16.5 STDEV : 10.5 VarCo : 0.584
2.4.16	Total : 27950 μ sec Mean : 931 μ sec Success rate: 100.00% Total iter: 451 Avg iter time: 62 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 451 Min : 1 Max : 49 Mean : 15.0 Median: 12.5 STDEV : 10.3 VarCo : 0.685	Min : 1 Max : 49 Mean : 15.0 Median: 12.5 STDEV : 10.3 VarCo : 0.685
2.4.17	Total : 38057 μ sec Mean : 1268 μ sec Success rate: 100.00% Total iter: 710 Avg iter time: 54 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 710 Min : 2 Max : 71 Mean : 23.7 Median: 18.0 STDEV : 17.2 VarCo : 0.726	Min : 2 Max : 71 Mean : 23.7 Median: 18.0 STDEV : 17.2 VarCo : 0.726

Again, the configurations with very small “p” values complete 30 search executions faster than the baseline experiment. However, most of the search runtimes are so close to each other that 1000 search executions are run against the best 2 experiments (measured by total iterations) 2.4.12 and 2.4.16.

Experiment	Execution Overview Overview	Restarts	Total Iterations	Iterations to solve
2.4.18	Total : 970976 μ sec Mean : 970 μ sec Success rate: 100.00% Total iter: 14802 Avg iter time: 66 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 14802 Min : 1 Max : 61 Mean : 14.8 Median: 13.0 STDEV : 8.6 VarCo : 0.578	Min : 1 Max : 61 Mean : 14.8 Median: 13.0 STDEV : 8.6 VarCo : 0.578
2.4.18	Total : 1045715 μ sec Mean : 1045 μ sec Success rate: 100.00% Total iter: 18919 Avg iter time: 55 μ sec	Sum : 0 Min : 0 Max : 0 Mean : 0 Median: 0.0 STDEV : 0.0	Sum : 18919 Min : 1 Max : 100 Mean : 18.9 Median: 14.0 STDEV : 15.2 VarCo : 0.803	Min : 1 Max : 100 Mean : 18.9 Median: 14.0 STDEV : 15.2 VarCo : 0.803

Like seen with experiments against uf20-01, a very small “p” value (here 0.001) and a tabu of 10% the number of variables in the problem (here 2) produces the most performant search.

Finally, let us visualize the history for the “unsatisfied clauses” counter for both uf20 instances. Both graphs visualize WalkSAT’s careful variable selection that prioritises variables with the lowest “negative gain”.



2.5 Conclusion

The previous experiments have given insights into how the operation of WalkSAT/SKC is affected by its parameters “wp” and “tl”. The parameters “restart” and “max_iterations” have been evaluated in section 1 already and the concept remains untouched.

WalkSAT completed the search operations significantly faster than GWSAT so that “restart” and “max_iterations” can be throttled much stronger than with GWSAT depending on the problem instance.

The experiments have shown that WalkSAT favours very small “p” values of 0.001 or 0.05 and focus on its lowest “negative gain” metric otherwise. Too large “p” values render the WalkSAT search more or less into a random search risking getting into states with more unsatisfied clauses.

Empiric results indicate that a good value for the tabu list length (or better how many iterations a flipped variable should be blocked from being selected again) is around 10%-15% of the number of variables in a given problem. This indication of course would need to be further analysed against additional and more complex problems.

While the general aim of WalkSAT is the intensification of the search by limiting the search space to unsatisfied clauses only and preferring lower negative gains during the variable selection, the introduction of the Tabu list provides diversification of the search process.

WalkSAT’s small neighbourhood for every iteration (1 unsatisfied clause) and the resulting quasi-constant time required per iteration makes it very scalable to larger problem instances.

2.6 Debug WalkSAT

To debug the operation of the GWSAT implementation set the DEBUG flag to TRUE in file 'Leske_183658_WalkSAT.py'.

This out is based on a baseline run against uf50-01.

```
python Leske_183658_WalkSAT.py uf50-01.cnf 1 1000 10 0.4 5
```

Variable selection based on negative gain 0:

```
Iteration                : 1
Tabu List                : [None, None, None, None, None, None, None, None, None, None, None, None, None, None,
None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None,
None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None,
None]
candidate_vars before tabu: [41, 34, 20]
candidate_vars after tabu: [41, 34, 20]
candidate_vars neg_gain  : [2, 0, 1]
selected variable       : 34
...
```

This output shows that for 3 candidate variables the negative gain has been calculated. Variable 34 has a negative gain of zero and is selected.

Variable selection based on lowest negative gain:

```
...
Iteration                : 4
Tabu List                 : [8, None, None, None, None, None, None, None, 7, None, None, None, None, None,
None, None, None, None, None, None, None, None, None, None, None, None, None, None, None,
None, 6, None, None, None, None, None, None, None, None, None, None, None, None, None, None, None]
candidate_vars before tabu: [25, 50, 18]
candidate_vars after  tabu: [25, 50, 18]
random walk
candidate_vars neg_gain   : [2, 2, 1]
selected variable         : 18
...
```

This output shows that none of the candidate variables has a negative gain of zero. Hence, the variable with the lowest negative gain is selected.

Variable selection based on random walk with probability “p”:

```
Iteration                : 14
Tabu List                : [8, None, 14, 11, None, None, 16, None, 7, 15, None, None, 17, None, None,
None, None, 9, None, None, None, None, None, None, None, None, None, 13, None, None, None, 12, None, 6,
None, None, None, 10, None, None, None, None, None, None, None, None, None, None, None, 18, None]
candidate_vars before tabu: [32, 26, 6]
candidate_vars after tabu: [32, 26, 6]
random walk
candidate_vars neg_gain  : [3, 2, 2]
selected variable       : 6
...
```

This output shows how parameter “p” fired and a random variable was selected. A side effect of the limited number of candidate variables with WalkSAT is the still the random walk step can be an optimal step for the algorithm. This observation is enforced, when 1 or 2 variables are tabued in a given iteration.

Maintaining tabu list:

```
...
Iteration          : 14
Tabu List           : [8, None, 14, 11, None, None, 16, None, 7, 15, None, None, 17, None, None,
None, None, 9, None, None, None, None, None, None, None, None, 13, None, None, None, 12, None, 6,
None, None, None, 10, None, None, None, None, None, None, None, None, None, None, 18, None]
candidate_vars before tabu: [32, 26, 6]
candidate_vars after tabu: [32, 26, 6]
random walk
candidate_vars neg_gain : [3, 2, 2]
selected variable      : 6

Iteration          : 15
Tabu List           : [8, None, 14, 11, None, 19, 16, None, 7, 15, None, None, 17, None, None, None,
None, 9, None, None, None, None, None, None, None, None, 13, None, None, None, 12, None, 6, None,
None, None, 10, None, None, None, None, None, None, None, None, None, None, 18, None]
candidate_vars before tabu: [16, 38, 6]
candidate_vars after tabu: [16, 38]
candidate_vars neg_gain : [0, 1]
selected variable      : 16
...
```

This output shows that in iteration 14 variable 6 was selected. The tabu list for iteration 17 shows that variable 6 is blocked until and including iteration 19. Starting iteration 20 variable 6 can be selected again.

Filtering candidates based on tabu list:

```
...
Iteration          : 258
Tabu List           : [223, 236, 217, 151, None, 230, 250, 247, 238, 262, 175, 243, 84, 227, 258,
159, 213, 176, 237, 259, 251, 210, 260, 222, 209, 189, 249, 171, 164, 242, 221, 196, 261, 123, 193, 245,
248, 231, None, 197, 244, 254, 143, 239, 257, 252, 195, 246, 103, 256]
candidate_vars before tabu: [33, 4, 12]
candidate_vars after tabu: [4, 12]
candidate_vars neg_gain : [5, 3]
selected variable      : 12

Iteration          : 259
Tabu List           : [223, 236, 217, 151, None, 230, 250, 247, 238, 262, 175, 263, 84, 227, 258,
159, 213, 176, 237, 259, 251, 210, 260, 222, 209, 189, 249, 171, 164, 242, 221, 196, 261, 123, 193, 245,
248, 231, None, 197, 244, 254, 143, 239, 257, 252, 195, 246, 103, 256]
candidate_vars before tabu: [24, 50, 5]
candidate_vars after tabu: [24, 50, 5]
candidate_vars neg_gain : [0, 2, 3]
selected variable      : 24

...
Iteration          : 263
Tabu List           : [223, 236, 217, 151, None, 230, 250, 247, 238, 262, 175, 263, 84, 227, 258,
159, 213, 176, 237, 266, 251, 210, 260, 264, 209, 189, 249, 171, 164, 242, 221, 196, 261, 123, 193, 245,
248, 231, None, 197, 244, 254, 143, 239, 257, 267, 195, 246, 265, 256]
candidate_vars before tabu: [32, 12, 46]
candidate_vars after tabu: [32]
candidate_vars neg_gain : [1]
selected variable      : 32
...
```

From this snippet we can see that variable 12 was selected in iteration 258. Therefore the debug for iteration 259 shows that variable needs to idle for 5 subsequent iterations (259, 260, 261, 262, 263) and can only be used again thereafter.

The debug for iteration 263 indeed confirms that variable 12 is still blocked by the tabu list and removed from the candidate list

3 Run-Time Distribution Evaluation

Performing Local Search for a SAT problem represents a decision problem in which a candidate variable assignment is determined to satisfy the given SAT instance or not. Depending on the randomness involved in creating the initial variable assignment and thereafter the decisions between “normal” and “random” search steps, the runtime for such local search algorithms experiences variance. Hence, the Local Search for SAT problems is defined by the Las Vegas class of algorithms where “runtime” is treated as the random variable that requires quantification.

The statistics collected in sections 1 and 2 already outline a few key findings:

- WalkSAT requires far less iterations and runtime to solve the uf20 or uf50 instances compared to GWSAT.
- Due to the limited search space per iteration WalkSAT offers a quasi-constant runtime per iteration.
- GWSAT runtime per iteration can be reduced by increasing the “wp” parameter at the cost of a more and more random search.
- Instance uf20-02 can be solved much faster both in runtime and iterations than uf20-01.

In order to evaluate the run-time distribution each problem instance uf20-01 and uf20-02 will be executed 1000 times against GWSAT and WalkSAT. Both instances will be run against each algorithm with the same parameters to allow comparison. The following parameters will be used for the run-time distribution evaluation:

Experiment	Algorithm	Instance	Executions	Restarts	Iterations	wp	tl
3.1.1	GWSAT	uf20-01	1000	3	200	0.4	n/a
3.1.2	GWSAT	uf20-02	1000	3	200	0.4	n/a
3.1.3	WalkSAT	uf20-01	1000	3	200	0.001	3
3.1.4	WalkSAT	uf20-02	1000	3	200	0.001	3

For each algorithm these parameters have proven to be good candidates for the complexity of the uf20 SAT instances.

	GWSAT uf20-01	GWSAT uf20-02	WalkSAT uf20-01	WalkSAT uf20-02
Overview	Total : 1005598 μ sec	Total : 5580940 μ sec	Total : 1610836 μ sec	Total : 979001 μ sec
	Mean : 10055 μ sec	Mean : 5580 μ sec	Mean : 1610 μ sec	Mean : 979 μ sec
	Success rate: 100.00%	Success rate: 100.00%	Success rate: 100.00%	Success rate: 100.00%
	Total iter: 54989	Total iter: 29299	Total iter: 30983	Total iter: 15780
	Avg iter time: 183 μ sec	Avg iter time: 190 μ sec	Avg iter time: 52 μ sec	Avg iter time: 62 μ sec
Runtime in μ sec	Sum : 10273975	Sum : 5580940	Sum : 1610836	Sum : 979001
	Min : 825	Min : 684	Min : 346	Min : 363
	Max : 63683	Max : 34548	Max : 7510	Max : 3243
	Mean : 10274	Mean : 5581	Mean : 1611	Mean : 979
	Median: 7416.5	Median: 4573.5	Median: 1303.0	Median: 863.0
	STDEV : 9170	STDEV : 3992	STDEV : 1093	STDEV : 430
	VarCo : 0.893	VarCo : 0.715	VarCo : 0.679	VarCo : 0.44
	Q0.1 : 2340	Q0.1 : 1793	Q0.1 : 623	Q0.1 : 560
	Q0.25 : 3972	Q0.25 : 2727	Q0.25 : 810	Q0.25 : 681
	Q0.5 : 7416	Q0.5 : 4573	Q0.5 : 1303	Q0.5 : 863
Total Iterations	Q0.75 : 13175	Q0.75 : 7157	Q0.75 : 2043	Q0.75 : 1159
	Q0.9 : 22857	Q0.9 : 10741	Q0.9 : 3109	Q0.9 : 1531
	Sum : 54989	Sum : 29299	Sum : 30983	Sum : 15780
	Min : 2	Min : 2	Min : 2	Min : 2
	Max : 379	Max : 181	Max : 176	Max : 67
	Mean : 55.0	Mean : 29.3	Mean : 31.0	Mean : 15.8
	Median: 39.0	Median: 23.5	Median: 24.0	Median: 13.0
	STDEV : 49.9	STDEV : 22.3	STDEV : 25.9	STDEV : 10.0
	VarCo : 0.908	VarCo : 0.762	VarCo : 0.837	VarCo : 0.631
	Q0.1 : 11	Q0.1 : 8	Q0.1 : 8	Q0.1 : 6
	Q0.25 : 21	Q0.25 : 13	Q0.25 : 12	Q0.25 : 9
	Q0.5 : 39	Q0.5 : 23	Q0.5 : 24	Q0.5 : 13
	Q0.75 : 71	Q0.75 : 37	Q0.75 : 40	Q0.75 : 20
	Q0.9 : 120	Q0.9 : 58	Q0.9 : 66	Q0.9 : 29

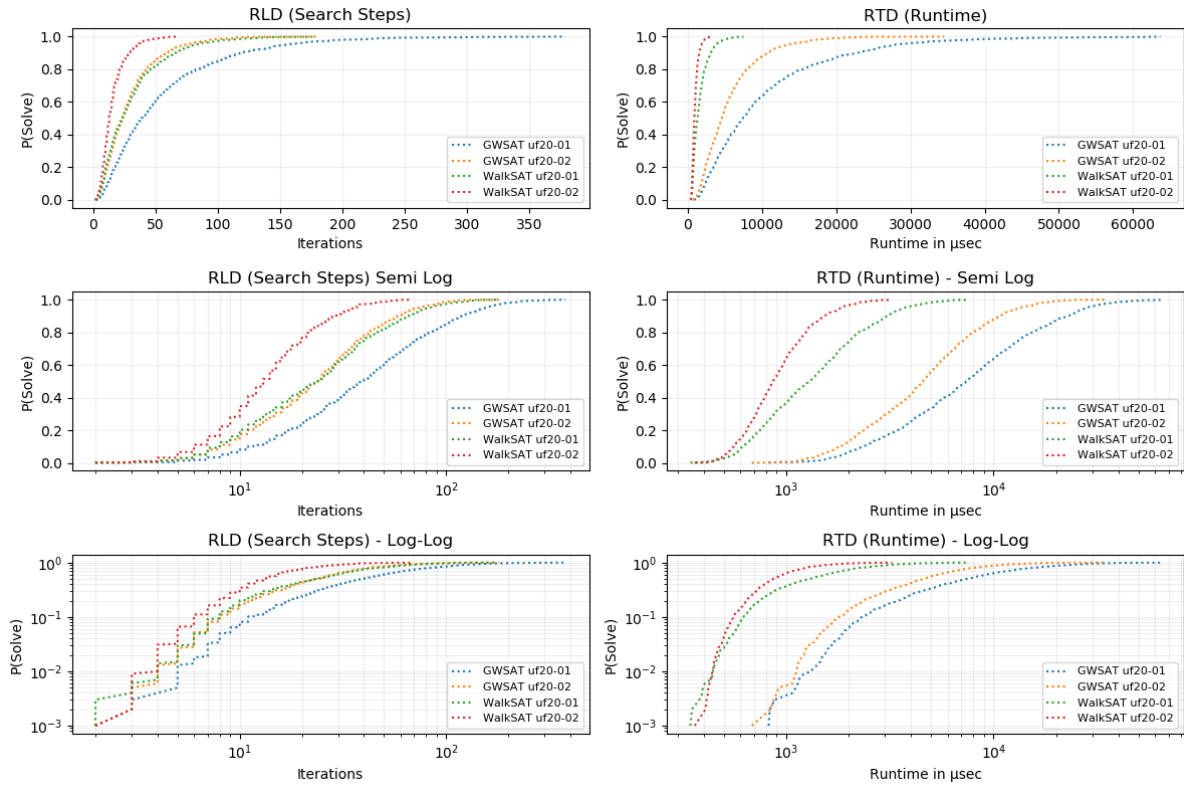


Figure 8 Run-Time Distribution for uf20-01 and uf20-02 against GWSAT and WalkSAT

Figure 8 visualizes many of the statistics given in the table above. Especially the run-time distribution graph clearly demonstrates that:

- Problem instances uf20-01 and uf20-02 are of different complexity levels
- WalkSAT solves both instances 6x to 7x faster than GWSAT

While WalkSAT solved ~75% uf20-01 and ~95% uf20-02 instances in less then 2000 μ sec, GWSAT reached those success rate levels only far beyond the 10000 μ sec runtime.

The Run-Length Distribution based on the Search Steps (or Iterations) to complete all executions highlights a very interesting detail. The number of search iterations GWSAT needs for solving instance uf20-02 matches the number of search iterations WalkSAT needs for solving instance uf20-01.

The statistics collected also reveal that each algorithm is able to solve 1000 executions of instance uf20-02 in ~50% of the search iterations required for solving 1000 executions of instance uf20-01. Similarly, instance uf20-01 can be solved 1000 times by WalkSAT in 50% the time GWSAT requires for the same task. The same statement can be made about instance uf20-02.

These findings again underline that the complexity for solving SAT instances of a given size can vary a lot – as do the 2 instances to be analysed.

The runtime statistics also show that a single WalkSAT iteration (52 & 62 μ sec) is approximately 3-4 times faster than a GWSAT iteration (183 & 190 μ sec). This can be explained with the reduced neighbourhood WalkSAT operates on (1 unsatisfied clause). As explained already in section 1, the GWSAT iterations runtime is largely affected by the random walk probability “wp”, which helps GWSAT avoiding analysing the outcome of flipping every variable by just randomly picking a variable to flip.

The cost for initially setting up the problem including reading the problem instance and filling data structures to hold the instance information (variables, literals, clauses) was measured to be ~320 μ sec.

```
search = LocalSearch(cnf_file, restarts, max_iterations, alg, p=p, tl=tl, DEBUG=DEBUG)
```

The start of a search execution – be it initial start or a restart – involves generating a random variable assignment and counting the number of unsatisfied clauses. This was measured to take ~170 μ sec for GWSAT.

```
self.cnf.initial_solution()
```

For WalkSAT we need to also account for the time to create and initialize the Tabu list. Hence the cost of starting or restarting a search is slightly higher at ~174 μ sec.

```
self.cnf.initial_solution()
```

```
if self.alg == 'walksat':  
    tabu = [ None for x in self.cnf.variables ]
```

Note: All initialization costs (in μ sec) given above are only relevant for the uf20 instances to be analysed in the assignment. Larger problem instances will result in appropriately larger setup costs.

One can conclude that we the given experiment setup the cost of a restart is approximately equal to:

- 1 GWSAT iteration
- 3 WalkSAT iterations

Of course, one also needs to add to the restart cost the time penalty required for the iterations to work down the number of unsatisfied clauses. Hence, the restart cost should be accounted much higher than the pure time penalty to create a new random variable assignment and counting the number of unsatisfied clauses initially. This explains why instances uf20-01 and uf20-02 are solved the fastest when avoiding search restarts (or keeping them to a bare minimum as in GWSAT for uf20-01).

The following table is a subset of information given from the above run-time distribution experiments and repeated here to provide context for the following calculations.

	GWSAT uf20-01	GWSAT uf20-02	WalkSAT uf20-01	WalkSAT uf20-02
Overview	Total : 10055998 μ sec	Total : 5580940 μ sec	Total : 1610836 μ sec	Total : 979001 μ sec
	Mean : 10055 μ sec	Mean : 5580 μ sec	Mean : 1610 μ sec	Mean : 979 μ sec
	Success rate: 100.00%	Success rate: 100.00%	Success rate: 100.00%	Success rate: 100.00%
	Total iter: 54989	Total iter: 29299	Total iter: 30983	Total iter: 15780
	Avg iter time: 183 μ sec	Avg iter time: 190 μ sec	Avg iter time: 52 μ sec	Avg iter time: 62 μ sec
Runtime in μsec	Q0.75 : 13175	Q0.75 : 7157	Q0.75 : 2043	Q0.75 : 1159
	Q0.9 : 22857	Q0.9 : 10741	Q0.9 : 3109	Q0.9 : 1531
Total Iterations	Max : 379	Max : 181	Max : 176	Max : 67
	Q0.75 : 71	Q0.75 : 37	Q0.75 : 40	Q0.75 : 20
	Q0.9 : 120	Q0.9 : 58	Q0.9 : 66	Q0.9 : 29

With the average time required per iteration and the setup costs mentioned above we can now define upper time limits to solve e.g. 75% or 90% of the uf20 problem instances. Accounting for the restart cost is only required for the max iterations example for GWSAT uf20-01. All other experiments solved all executions without restart.

GWSAT:

uf20-01

75%: 320 μ sec + 71*183 μ sec = 13313 μ sec or 13.3 msec
90%: 320 μ sec + 120*183 μ sec = 22280 μ sec or 22.3 msec
Max: 320 μ sec + 379*183 μ sec + 170 μ sec = 69847 μ sec or 69.8 msec

uf20-02

75%: 320 μ sec + 37*190 μ sec = 7350 μ sec or 7.4 msec
90%: 320 μ sec + 58*190 μ sec = 11340 μ sec or 11.3 msec
Max: 320 μ sec + 67*190 μ sec = 13050 μ sec or 13.1 msec

WalkSAT:

uf20-01

75%: 320 μ sec + 40*52 μ sec = 2400 μ sec or 2.4 msec
90%: 320 μ sec + 66*52 μ sec = 3752 μ sec or 3.8 msec
Max: 320 μ sec + 176*52 μ sec = 9472 μ sec or 9.5 msec

uf20-02

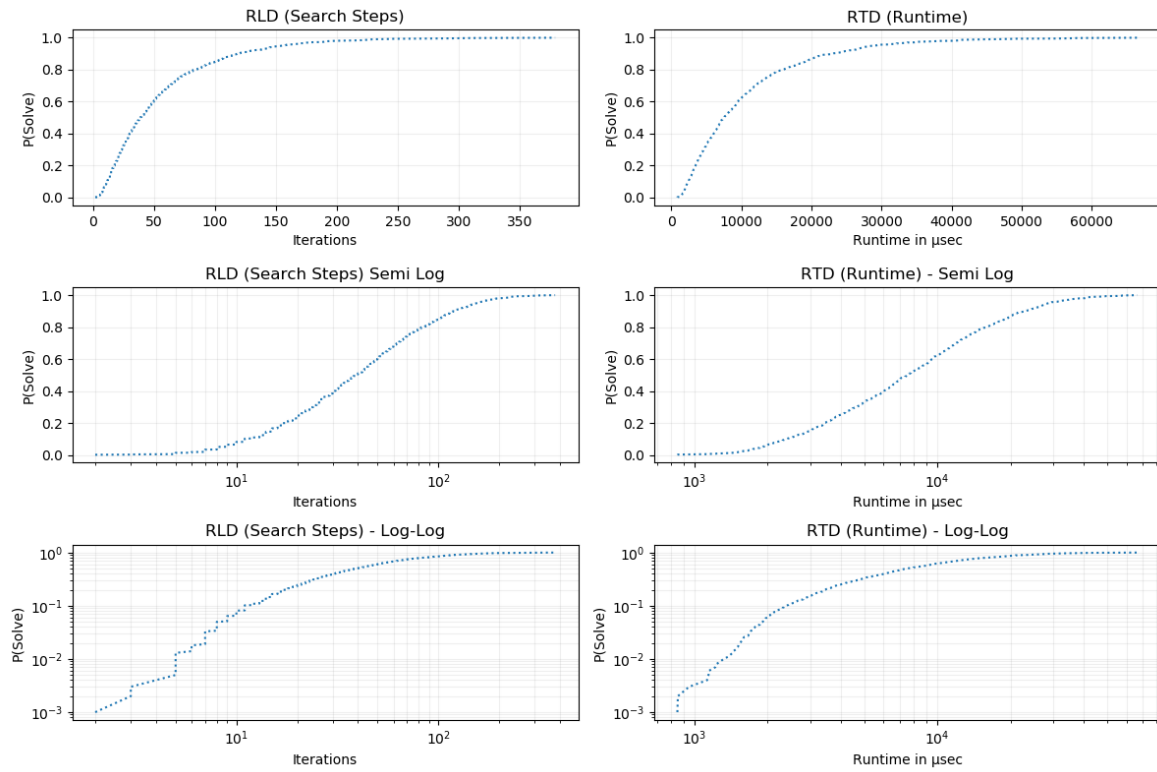
75%: 320 μ sec + 20*62 μ sec = 1560 μ sec or 1.6 msec
90%: 320 μ sec + 29*62 μ sec = 2118 μ sec or 2.1 msec
Max: 320 μ sec + 66*62 μ sec = 4412 μ sec or 4.4 msec

The calculated runtime based on setup time plus runtime per iteration matches roughly the runtime statistics taken for the experiments (which do not contain the setup time). Especially for instance uf20-01 it is obvious that 90% of the executions a solution is found in $\sim 1/3$ of the time required to find a solution in 100% of the executions.

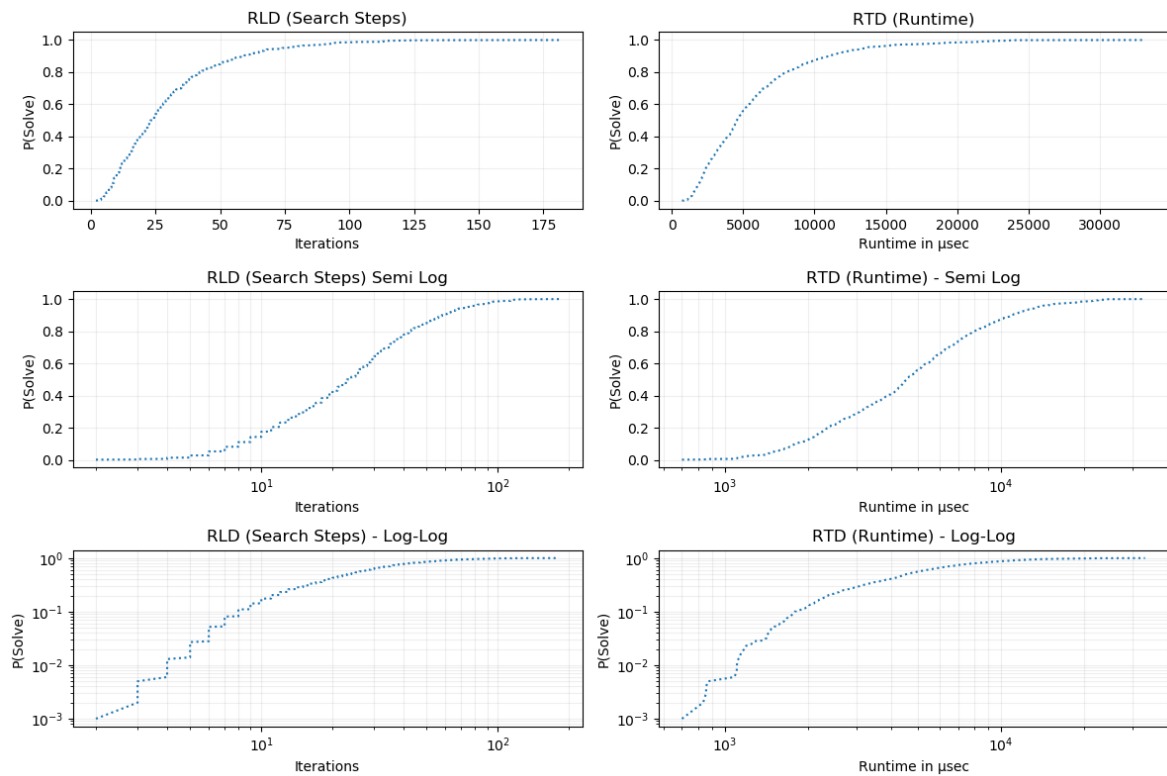
Appendix A

This appendix provides RTD diagrams for the individual RTD experiments for completeness.

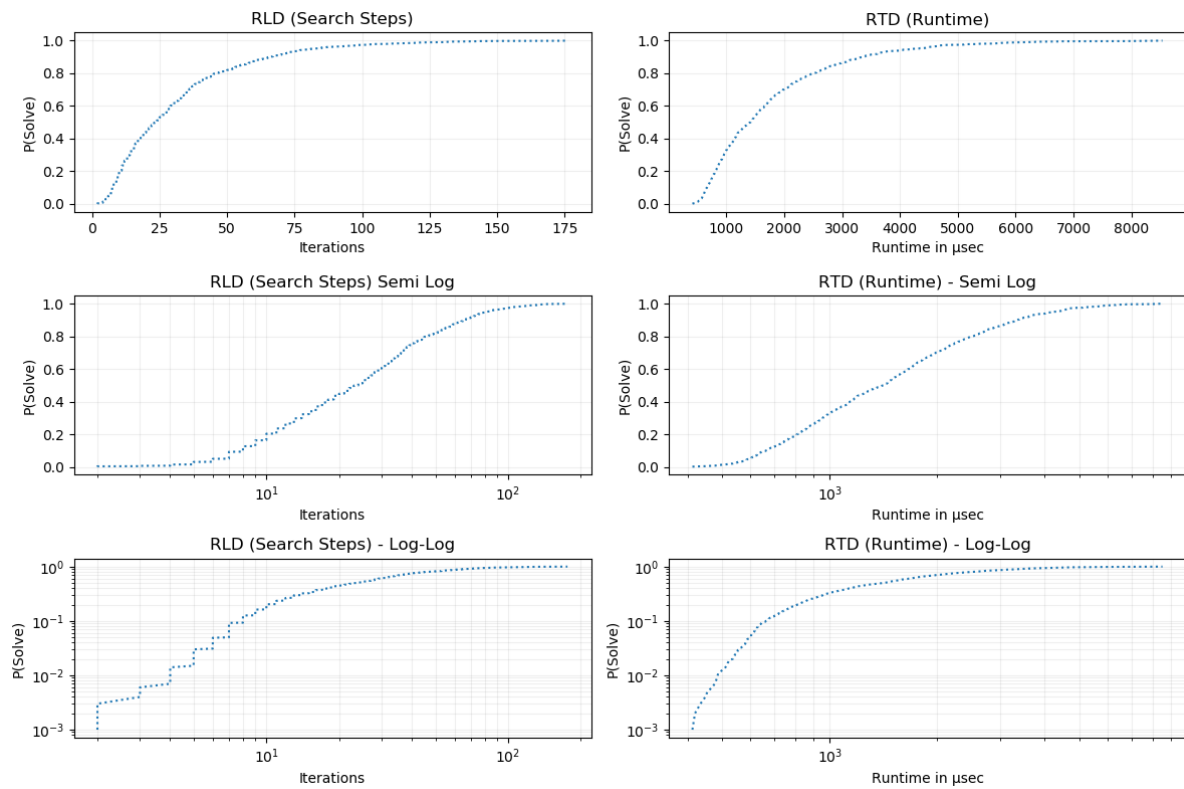
GWSAT uf20-01:



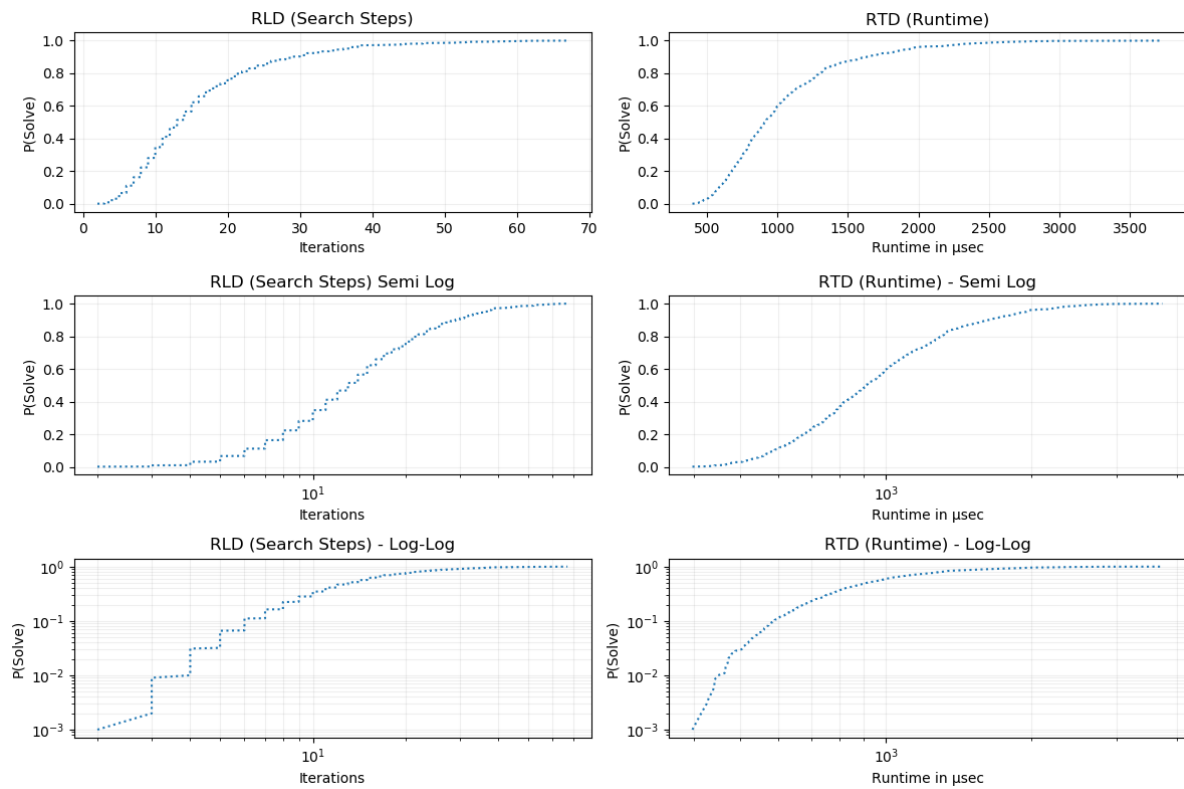
GWSAT uf20-02:



WalkSAT uf20-01:



WalkSAT uf20-02:



References

- [1] <https://pymotw.com/3/time/>